



Simpsons Character Recognition Project

Siqi Li, Lirong Ma, Zhangyi (Rocky) Ye, Haonan Zhang,
Yuxuan (Nancy) Zhang, Luping (Rachel) Zhao



The Simpsons

Popular American
animated sitcom
On air since 1989



Problem Statement

Problem: Large number of characters, sometimes we don't know who characters are while watching the show.

Goal: build deep learning models to detect and classify Simpsons characters.

Deployment:

Real-time character detection and classification.

Viewers know who they are watching without pressing pause to check their phones.

Dataset

Public dataset

18,992 pictures for 18 characters

Divide into 3 sets of data

Training (60%)

Validation (20%)

Test (20%)

Dataset Examples - Easy



Dataset Examples - Medium



Dataset Examples - Hard



Dataset - Challenge

Wrong character labels



Missing bounding box labels for images with multiple characters

Affect our ability to test model performance for part 2 detection





Technical Approach

Data Cleaning & Augmentation

Convert images to pixels and normalized

Randomly rotate images (rotationrange=15)

Randomly zoom inside pictures (zoom_range=0.2)

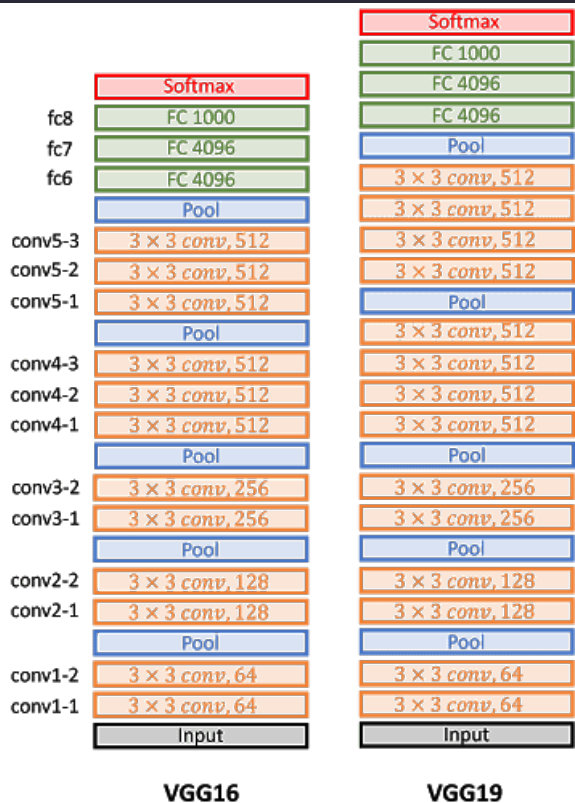
Randomly apply shearing transformations (shear_range=0.2)

Randomly flip images horizontally

Randomly shift images horizontally/vertically (shift_range=0.2)



VGG16/VGG19



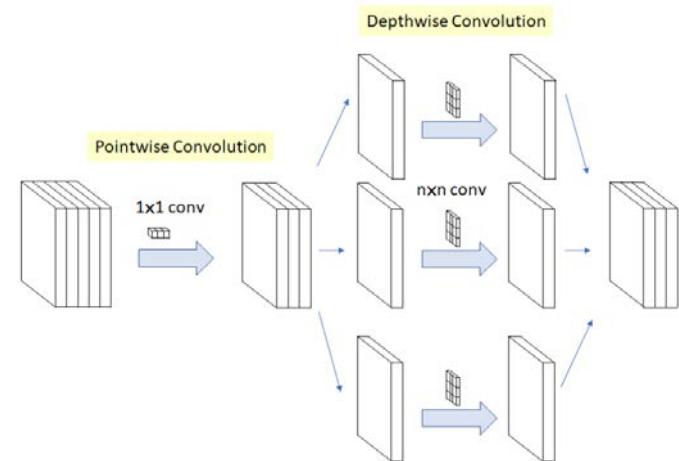
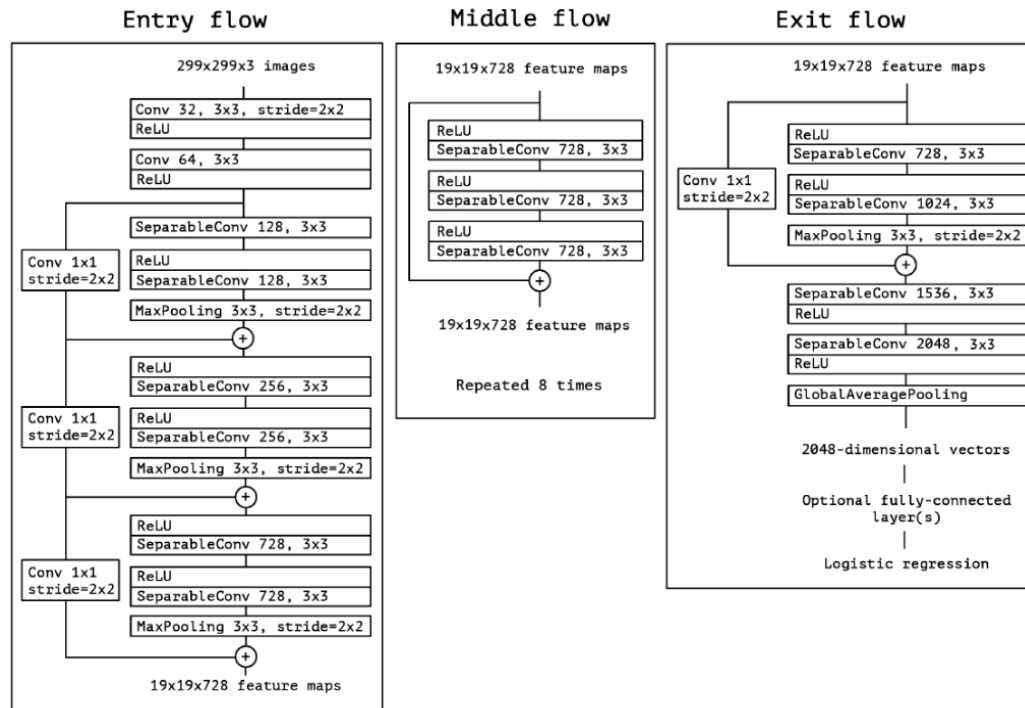
VGG16

- Convolutional layers
- Pooling layer
- Flatten layer
- Dropout layer

VGG19

- Similar to VGG16, except it has four convolutional layers in the fourth and fifth block.

Xception (Extreme version of Inception)




Object Detection: Faster R-CNN

Use sliding window and apply a CNN to many different crops of the image?

Inflexible Size


Too computationally expensive



Fixed Region Proposals?


Select 2000 region proposals in a few seconds on CPU and apply a CNN to each one of them?

Better, but not enough! The region proposals should be learned




Object Detection: Faster R-CNN

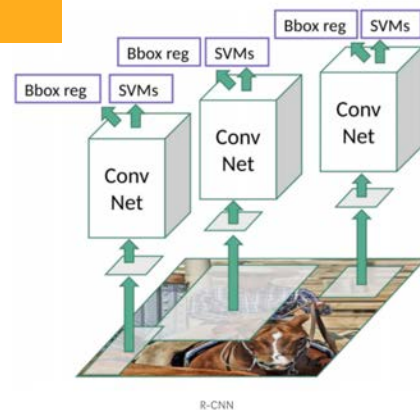
Faster R-CNN



Fast R-CNN



R-CNN



Runtime of the same task:

R-CNN: 49 hours

Fast R-CNN: 2.3 hours

Faster R-CNN: 0.2 hours

Part I – Simpson Characters Classification

Evaluation Metrics

Test Accuracy

Test Loss

Precision

Recall



VGG 16




Image size (32,32)




Image size (128,128)

VGG 19

Training vgg19 blocks=5, conv=15, dense=1, fsize=128, epoch=100




Image size (32,32)

Training vgg19 blocks=5, conv=15, dense=1, fsize=128, epoch=100

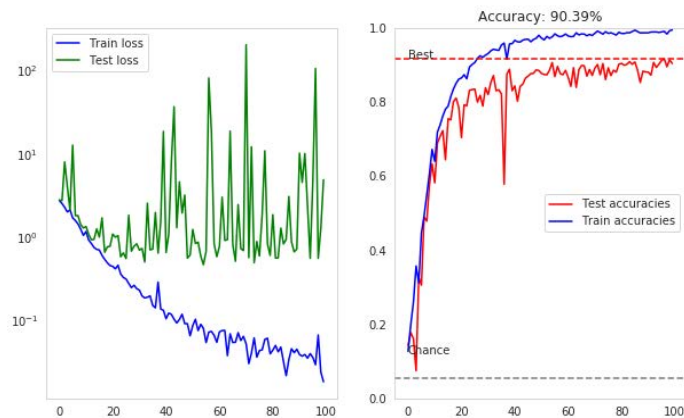


Image size (128,128)

Xception: Experiments

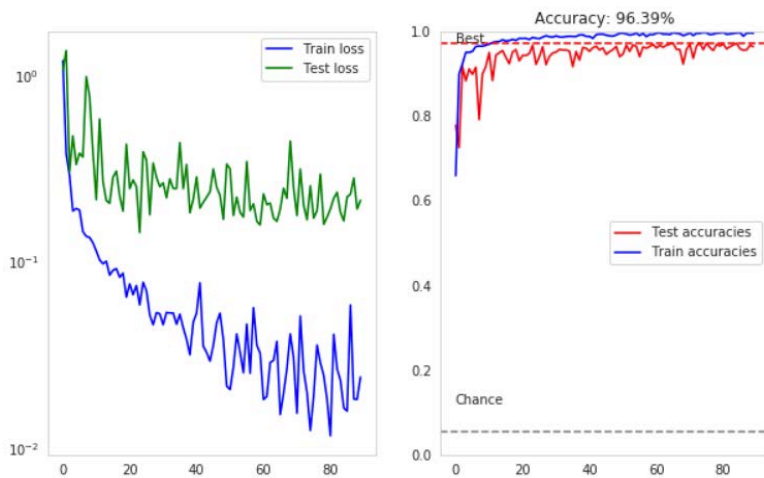


Image size (128,128)
No augmentation




Image size (256,256)
With augmentation

Best Model: Loss Curve

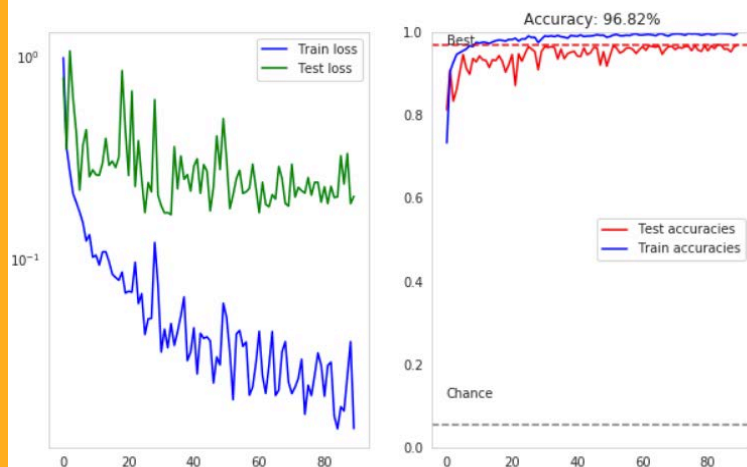
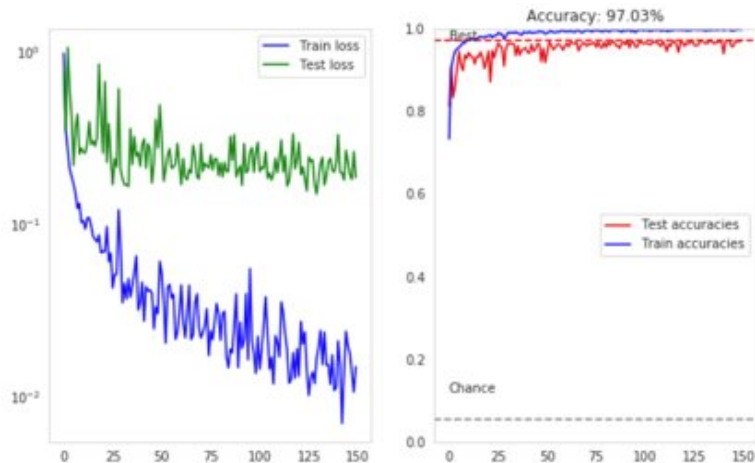


Image size: (128,128)

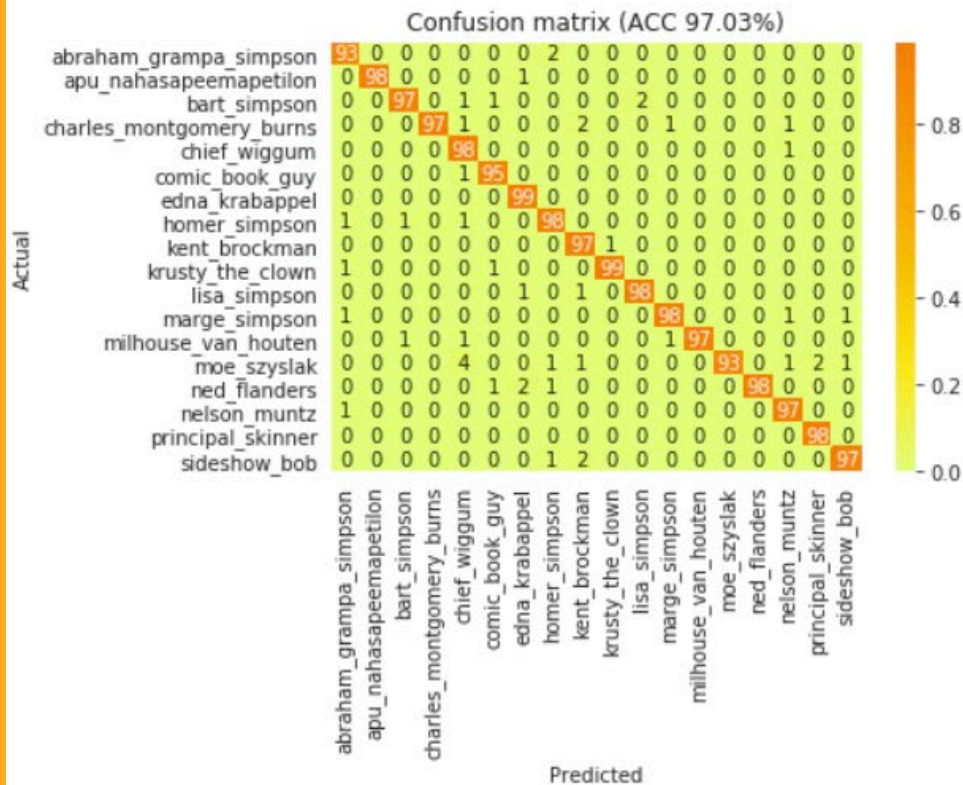
Augmentation:

- randomly shift images horizontally and vertically
- flip images horizontally




Test loss curve stabilized after ~50 epochs and test accuracy achieved 97.03% after 150 epochs

Best Model: Confusion Matrix



Best Model: Heatmap



Correct Classification

Misclassification

Most of the time, the model focusing on the correct part of the image

The model was unable to classify correctly when dissecting background

Part II - Object Detection and Classification

Evaluation metrics

Model level

Accuracy score = number of accurately classified characters /
actual number of bounding boxes

Character level

Precision: fraction of relevant characters among the retrieved
instances


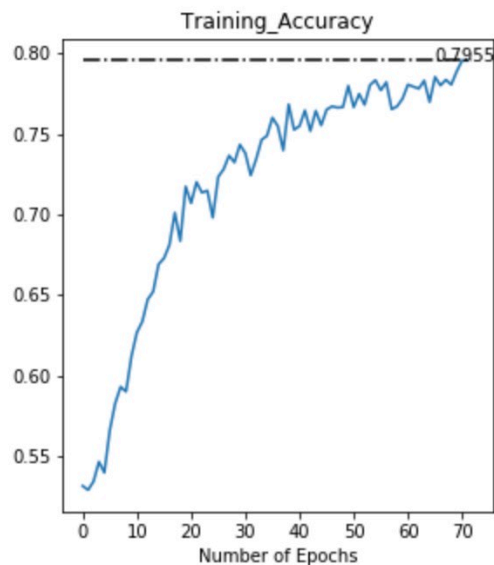

$\text{true positive cases} / \text{predicted positive cases}$

Recall: fraction of total amount of relevant characters that were
actually retrieved

$\text{true positive cases} / \text{actual positive cases}$

F1-score: a special weighted average of precision and recall

Best Model: Fast R-CNN



Part II Output Summary



Test Metrics Performance - Top Six Accurate Characters

Character	Precision	Recall	F1 score
ned_flanders	0.913	0.929	0.921
marge_simpson	0.832	0.966	0.894
kent_brockman	0.800	1.000	0.889
principal_skinner	0.870	0.895	0.883
krusty_the_clown	0.826	0.927	0.874
chief_wiggum	0.939	0.795	0.861

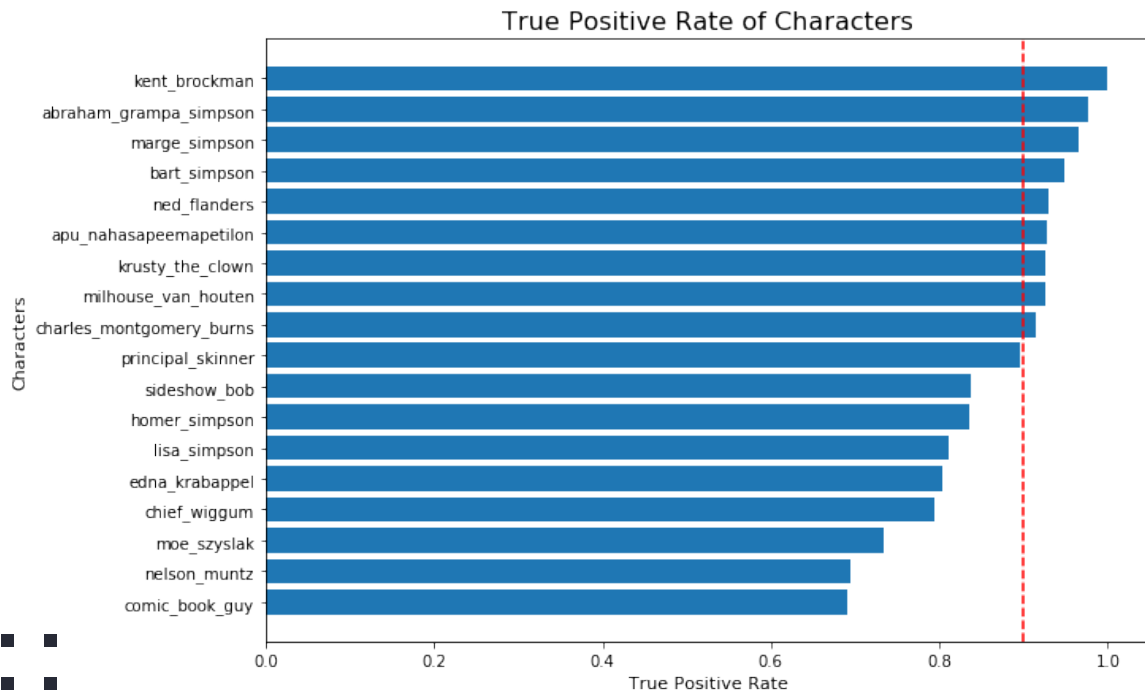
Part II Output Summary



Test Metrics Performance - Bottom Six Accurate Characters

Character	Precision	Recall	F1 score
abraham_grampa_simpson	0.636	0.977	0.771
charles_montgomery_burns	0.637	0.914	0.751
nelson_muntz	0.806	0.694	0.746
bart_simpson	0.541	0.949	0.689
moe_szyslak	0.629	0.733	0.677
comic_book_guy	0.547	0.690	0.611

Part II Output Summary

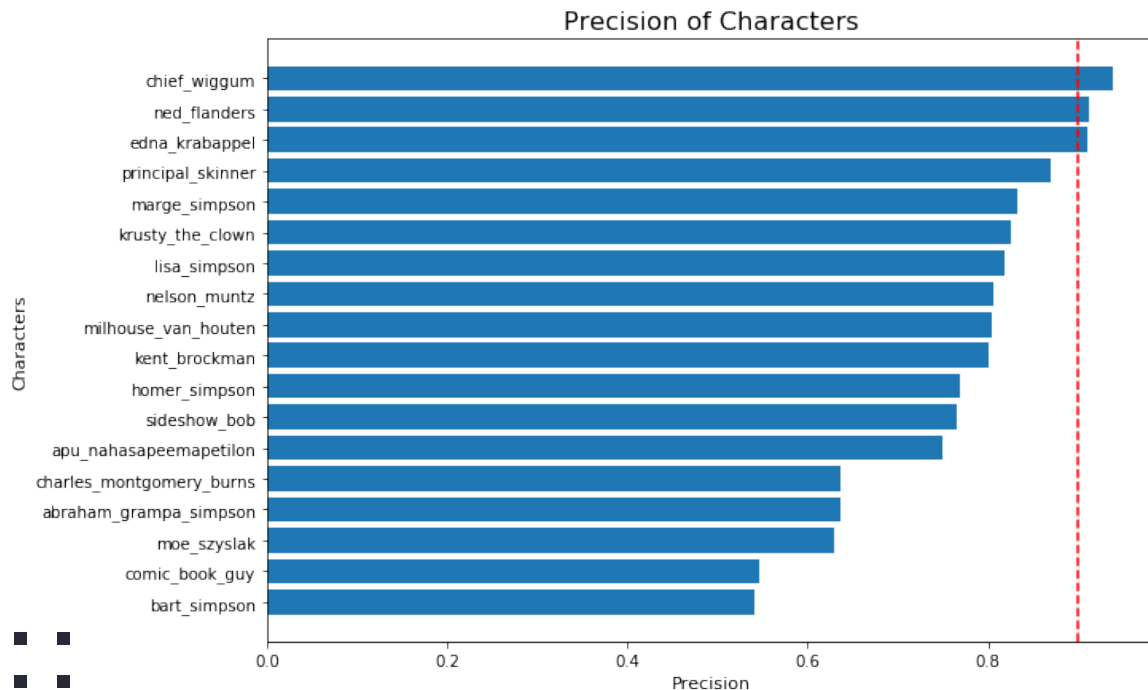


Half of the characters have recall greater than 90%.

The characters from Simpson Family (Marge, Bart, etc.) are highly likely to be identified.

The ranks of Simpson characters do not show significant difference.

Part II Output Summary





Only three characters have precision greater than 90%.



Some major characters (Bart Simpson) only have precision below 60%.

The model tends to predict more boxes as the major characters such as Bart Simpson and Abraham Grampa Simpson.

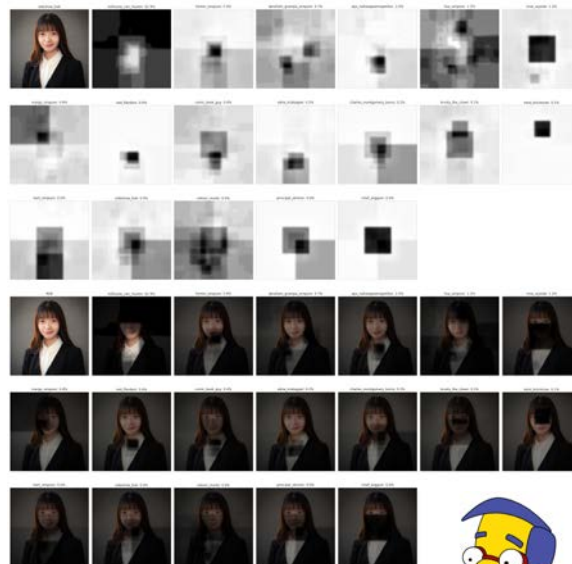

Easter Egg



Easter Egg



Easter Egg



Conclusion - Classification

Great performance for classification (Part I)

Misclassification due to multiple characters in the same image

CNN algorithms play an important role.



Conclusion - Detection

Good performance for detection (Part II)

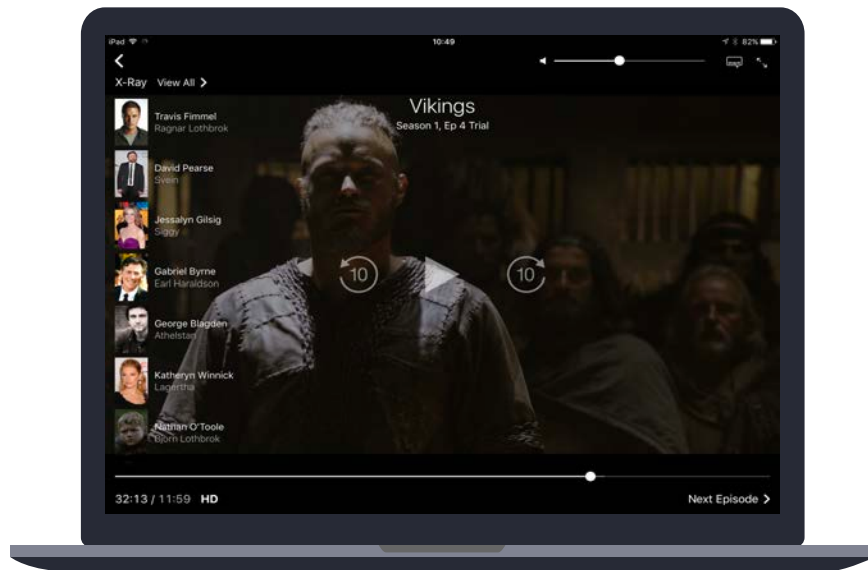
Errors due to missing bounding box labels in the pictures.

Training and predicting takes incredibly long





Application

Real time TV show
character detection
for people not good at
remembering names



Future Opportunities


Faster Algorithm: YOLO (v4)



References

- Attia, A. (2017, June 12). The Simpsons characters recognition and detection using Keras (Part 1). Retrieved from <https://medium.com/alex-attia-blog/the-simpsons-character-recognition-using-keras-d8e1796eae36>
- Carremans, B (2018, August 17). Classify butterfly images with deep learning in Keras. Retrieved from <https://towardsdatascience.com/classify-butterfly-images-with-deep-learning-in-keras-b3101fe0f98>
- He, K et al. (2015). Deep Residual Learning for Image Recognition. Retrieved from <https://arxiv.org/pdf/1512.03385.pdf>
- Krizhevsky, A et al. (2012). ImageNet Classification with Deep Convolutional Neural Networks. Retrieved from <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>
- Simonyan, K and Zisserman, A (2015). Very Deep Convolutional Networks for Large-Scale Image recognition. Retrieved from <https://arxiv.org/pdf/1409.1556.pdf>
- Szegedy, C et al. (2014). Going deeper with convolutions. Retrieved from <https://arxiv.org/abs/1409.4842>

the SIMPSONS



Thanks!

Any questions?



When you move on to
Deep Learning

